

CS2610 Lab 5 Report

Performance profiling with perf

Contents

- ❖ Summary of the blocking algorithm
- ❖ Observations
- ❖ Tables (values)
 - Matrix size: 128
 - Matrix size: 256
 - Matrix size: 512

Summary of the blocking algorithm

The cache blocking algorithm for matrix multiplication takes advantage of locality principles by multiplying blocks of matrices instead of the matrices as a whole. The matrices are split into blocks of different orders, which are powers of 2 (say, 4×4 , 16×16 , etc., or generally $b \times b$), so we will have, for each matrix, $(n/b)^2$ blocks. The algorithm can be described in two layers: block multiplication and element multiplication. The block multiplication layer (the outer 3 loops in the C++ code) selects the blocks from matrices **a** and **b** to multiply, simultaneously deciding which block in **c** to store the values to. This is effectively the same as regular matrix multiplication, if we treat our blocks as elements of a regular matrix. The element multiplication layer does the actual element-to-element multiplication across these blocks. For every pair of multipliable blocks we choose from **a** and **b**, we pick multipliable pairs of elements (in the three internal loops) and multiply them and increment the appropriate element of **c** by the obtained value. This procedure exploits locality within the L1 cache as long as each currently-used block from **a**, **b**, **c** remains in the cache throughout the block-wise multiplication. Each block occupies $4b^2$ bytes of memory, hence we require $12b^2$ to be less than the L1 cache size (32 KB for the machine on which the code was executed). Hence we can expect b upto 32 or 64 to show visible improvement.

Observations

Upon analyzing the output obtained using perf (included in a text file within the ZIP file, summarized in 3 tables following the observations), the following were observed:

For matrix size 128, the total number of cache references and cache misses remained nearly the same (with a pretty large increase in misses for block size 64), but the L1-dcache miss rate reduced significantly using the blocked algorithm, reaching its minimum for block sizes of 16 and 32 (local miss rate of 0.16%).

For matrix sizes 256 and 512, a good deal of improvement is visible with lower block sizes (≤ 16) wrt L1-dcache miss rate, and furthermore, a huge reduction in cache references is also seen in the blocked algorithm against the naive one! This detail showing up here but not in 128-sized matrices could be attributed to the fact that the matrices **a**, **b**, **c** can almost be accommodated entirely in the L1 cache when the size of the matrices is 128 (memory required for each matrix = $(2^7)^2 \cdot 4 = 64$ KB, whereas L1 cache size is 32 KB), which would prevent the naive algorithm from performing too badly since it can implement a good deal of locality by itself.

Considering L1-dcache misses as our benchmarking standard, ideal block size turns out to be any size ≤ 16 (though 16 doesn't perform as well as lower powers of 2 for 512, it is still a reasonable improvement over the naive algorithm), which matches closely with the prediction made at the end of the algorithm description.

Tables (values)

Matrix size: 128

Block size & Metric	0	4	8	16	32	64
Cache misses	58,689	42,833	51,650	66,966	56,112	89,941
Cache references	1,95,518	1,98,958	1,80,558	1,97,219	1,98,599	2,15,002
Cache miss rate	30.017%	21.529%	28.606%	33.955%	28.254%	41.833%
L1-dcache load misses	21,55,370	1,86,548	1,00,565	85,262	83,575	3,99,334
L1-dcache loads	4,86,72,059	5,61,88,622	5,30,77,639	5,17,99,811	5,12,18,394	5,09,46,732
L1-dcache stores	80,25,210	93,68,908	86,79,772	83,07,071	79,23,726	81,01,450
L1-dcache miss rate	4.43%	0.33%	0.19%	0.16%	0.16%	0.78%
L1-icache load misses	3,07,996	2,61,663	4,10,344	3,16,519	3,25,820	3,41,295

Matrix size: 256

Block size & Metric	0	4	8	16	32	64
Cache misses	1,15,680	1,18,507	80,955	1,09,994	89,003	66,634
Cache references	2,97,57,967	23,18,023	10,84,441	7,66,427	8,19,260	9,90,884
Cache miss rate	0.389%	5.112%	7.465%	14.352%	10.864%	6.725%
L1-dcache load misses	1,59,00,096	11,78,658	4,78,142	4,61,623	32,61,743	1,57,62,228
L1-dcache loads	34,86,53,939	40,24,79,679	37,25,63,672	36,42,87,031	36,01,53,918	36,36,21,853
L1-dcache stores	4,78,70,891	6,17,62,511	5,98,86,066	5,75,65,762	5,59,34,170	5,67,94,334
L1-dcache miss rate	4.56%	0.29%	0.13%	0.13%	0.91%	4.33%
L1-icache load misses	13,225	3,17,885	4,11,734	4,31,633	4,23,145	3,82,885

Matrix size: 512

Block size & Metric	0	4	8	16	32	64
Cache misses	4,13,693	4,36,642	4,19,514	2,43,635	3,70,614	3,66,016
Cache references	56,74,84,586	3,63,90,960	1,21,55,628	2,11,57,968	1,66,54,505	1,94,19,816
Cache miss rate	0.073%	1.200%	3.451%	1.152%	2.225%	1.885%
L1-dcache load misses	13,28,31,216	89,98,476	55,83,040	4,02,33,488	12,61,67,610	12,67,01,121
L1-dcache loads	2,67,08,29,600	3,11,00,23,247	2,90,15,26,102	2,81,29,50,177	2,81,79,86,947	2,77,10,44,535
L1-dcache stores	36,45,31,492	43,86,84,441	37,98,75,779	39,12,92,899	35,46,63,886	36,62,76,647
L1-dcache miss rate	4.97%	0.29%	0.19%	1.43%	4.48%	4.57%
L1-icache load misses	7,11,275	2,73,861	3,92,580	9,78,098	5,08,268	5,43,849